

data_wrangling

December 20, 2018

Link

Data Analysis with Python

1 Module 2 : Data Wrangling

1.0.1 Welcome!

By the end of this notebook, you will have learned the basics of Data Wrangling!

1.1 Table of contents

Identify and handle missing values

- Identify missing values
- Deal with missing values
- Correct data format

Data standardization

Data Normalization (centring/scaling)

Binning

Indicator variable

Estimated Time Needed: 30 min

1.2 What is the purpose of Data Wrangling?

Data Wrangling is the process of converting data from the initial format to a format that may be better for analysis.

1.2.1 What is the fuel consumption (L/100k) rate for the diesel car?

1.2.2 Import data

You can find the "Automobile Data Set" from the following link: <https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>. We will be using this data set throughout this course.

Import pandas

```
In [1]: import pandas as pd
```

1.3 Reading the data set from the URL and adding the related headers.

URL of dataset

```
In [2]: filename = 'https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.csv'
        print('URL read and saved as "filename"')
```

URL read and saved as "filename"

Python list "headers" containing name of headers

```
In [3]: headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration", "num-of-doors",
                  "drive-wheels", "engine-location", "wheel-base", "length", "width", "height", "curb-weight",
                  "num-of-cylinders", "engine-size", "fuel-system", "bore", "stroke", "compression-ratio",
                  "peak-rpm", "city-mpg", "highway-mpg", "price"]
```

Use the Pandas method `read_csv()` to load the data from the web address. Set the parameter "names" equal to the Python list "headers".

```
In [4]: df = pd.read_csv(filename, names = headers)
        print("Done")
```

Done

Use the method `head()` to display the first five rows of the dataframe.

```
In [5]: # To see what the data set looks like, we'll use the head() method.
        df.head()
```

```
Out[5]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	3	?	alfa-romero	gas	std	two	
1	3	?	alfa-romero	gas	std	two	
2	1	?	alfa-romero	gas	std	two	
3	2	164	audi	gas	std	four	
4	2	164	audi	gas	std	four	

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	

```

4          mpfi  3.19    3.40                8.0        115        5500        18

      highway-mpg  price
0           27  13495
1           27  16500
2           26  16500
3           30  13950
4           22  17450

[5 rows x 26 columns]

```

As we can see, several question marks appeared in the dataframe; those are missing values which may hinder our further analysis.

So, how do we identify all those missing values and deal with them?

How to work with missing data:

Steps for working with missing data: 1. Identify missing data 2. Deal with missing data 3. Correct data format

1. Identify and handle missing values

Convert "?" to NaN In the car dataset, missing data comes with the question mark "?". We replace "?" with NaN (Not a Number), which is Python's default missing value marker, for reasons of computational speed and convenience. Here we use the function:

to replace A by B

```
In [6]: import numpy as np
```

```

# replace "?" to NaN
df.replace("?", np.nan, inplace = True)
df.head(5)

```

```

Out[6]:   symboling  normalized-losses      make fuel-type aspiration num-of-doors \
0         3             NaN  alfa-romero    gas      std         two
1         3             NaN  alfa-romero    gas      std         two
2         1             NaN  alfa-romero    gas      std         two
3         2           164      audi      gas      std         four
4         2           164      audi      gas      std         four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size \
0  convertible      rwd      front      88.6  ...      130
1  convertible      rwd      front      88.6  ...      130
2   hatchback      rwd      front      94.5  ...      152
3      sedan      fwd      front      99.8  ...      109
4      sedan      4wd      front      99.4  ...      136

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg \
0          mpfi  3.47    2.68           9.0        111    5000    21
1          mpfi  3.47    2.68           9.0        111    5000    21
2          mpfi  2.68    3.47           9.0        154    5000    19
3          mpfi  3.19    3.40          10.0        102    5500    24

```

4	mpfi	3.19	3.40	8.0	115	5500	18
---	------	------	------	-----	-----	------	----

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

1.3.1 Evaluating for Missing Data

The missing values are converted to Python's default. We use Python's built-in functions to identify these missing values. There are two methods to detect missing data: 1. `.isnull()` 2. `.notnull()`

The output is a boolean value indicating whether the dataframe is missing data.

```
In [7]: missing_data = df.isnull()
missing_data.head(5)
```

```
Out[7]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	False	True	False	False	False	False	
1	False	True	False	False	False	False	
2	False	True	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	False	False	False	False	...	False	
1	False	False	False	False	...	False	
2	False	False	False	False	...	False	
3	False	False	False	False	...	False	
4	False	False	False	False	...	False	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

	city-mpg	highway-mpg	price
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False

[5 rows x 26 columns]

"True" stands for missing value, while "False" stands for not missing value.

1.3.2 Count missing values in each column

Using a for loop in Python, we can quickly figure out the number of missing values in each column. As mentioned above, "True" represents a missing value, "False" means the value is present in the dataset. In the body of the for loop the method ".value_counts()" counts the number of "True" values.

```
In [8]: for column in missing_data.columns.values.tolist():
        print(column)
        print (missing_data[column].value_counts())
        print("")
```

```
symboling
False      205
Name: symboling, dtype: int64
```

```
normalized-losses
False      164
True        41
Name: normalized-losses, dtype: int64
```

```
make
False      205
Name: make, dtype: int64
```

```
fuel-type
False      205
Name: fuel-type, dtype: int64
```

```
aspiration
False      205
Name: aspiration, dtype: int64
```

```
num-of-doors
False      203
True         2
Name: num-of-doors, dtype: int64
```

```
body-style
False      205
Name: body-style, dtype: int64
```

```
drive-wheels
False      205
Name: drive-wheels, dtype: int64
```

engine-location
False 205
Name: engine-location, dtype: int64

wheel-base
False 205
Name: wheel-base, dtype: int64

length
False 205
Name: length, dtype: int64

width
False 205
Name: width, dtype: int64

height
False 205
Name: height, dtype: int64

curb-weight
False 205
Name: curb-weight, dtype: int64

engine-type
False 205
Name: engine-type, dtype: int64

num-of-cylinders
False 205
Name: num-of-cylinders, dtype: int64

engine-size
False 205
Name: engine-size, dtype: int64

fuel-system
False 205
Name: fuel-system, dtype: int64

bore
False 201
True 4
Name: bore, dtype: int64

stroke
False 201
True 4

```

Name: stroke, dtype: int64

compression-ratio
False      205
Name: compression-ratio, dtype: int64

horsepower
False      203
True        2
Name: horsepower, dtype: int64

peak-rpm
False      203
True        2
Name: peak-rpm, dtype: int64

city-mpg
False      205
Name: city-mpg, dtype: int64

highway-mpg
False      205
Name: highway-mpg, dtype: int64

price
False      201
True        4
Name: price, dtype: int64

```

Based on the summary above, each column has 205 rows of data, with seven columns containing missing data:

1. "normalized-losses": 41 missing data
2. "num-of-doors": 2 missing data
3. "bore": 4 missing data
4. "stroke": 4 missing data
5. "horsepower": 2 missing data
6. "peak-rpm": 2 missing data
7. "price": 4 missing data

Deal with missing data How to deal with missing data:

1. Drop data
 - a. drop the whole row
 - b. drop the whole column
2. Replace data
 - a. replace it by mean

- b. replace it by frequency
- c. replace it based on other functions

Whole columns should be dropped only if most entries in the column are empty.

In our dataset, none of the columns are empty enough to drop entirely.

We have some freedom in choosing which method to replace data; however, some methods may seem more reasonable than others. We will apply each method to many different columns:

Replace by mean:

```
"normalized-losses": 41 missing data, replace them with mean
"stroke": 4 missing data, replace them with mean
"bore": 4 missing data, replace them with mean
"horsepower": 2 missing data, replace them with mean
"peak-rpm": 2 missing data, replace them with mean
```

Replace by frequency:

```
"num-of-doors": 2 missing data, replace them with "four".
* Reason: 84% sedans is four doors. Since four doors is most frequent, it is most likely to
```

Drop the whole row:

```
"price": 4 missing data, simply delete the whole row
* Reason: price is what we want to predict. Any data entry without price data cannot be used
```

Calculate the average of the column:

```
In [9]: avg_1 = df["normalized-losses"].astype("float").mean(axis = 0)
```

Replace "NaN" by mean value in "normalized-losses" column:

```
In [10]: df["normalized-losses"].replace(np.nan, avg_1, inplace = True)
```

Calculate the mean value for 'bore' column:

```
In [11]: avg_2=df['bore'].astype('float').mean(axis=0)
```

Replace NaN by mean value:

```
In [12]: df['bore'].replace(np.nan, avg_2, inplace= True)
```

Question #1:

According to the example above, replace NaN in "stroke" column by mean:

```
In [13]: avg3=df['stroke'].astype('float').mean(axis=0)
         df['stroke'].replace(np.nan, avg3, inplace=True)
         df['stroke']
```



```
Out[13]: 0      2.68
         1      2.68
         2      3.47
         3      3.40
         4      3.40
         5      3.40
         6      3.40
         7      3.40
         8      3.40
         9      3.40
        10      2.80
        11      2.80
        12      3.19
        13      3.19
        14      3.19
        15      3.39
        16      3.39
        17      3.39
        18      3.03
        19      3.11
        20      3.11
        21      3.23
        22      3.23
        23      3.39
        24      3.23
        25      3.23
        26      3.23
        27      3.39
        28      3.46
        29      3.90
         ...
        175     3.54
        176     3.54
        177     3.54
        178     3.35
        179     3.35
        180     3.35
        181     3.35
        182     3.40
        183     3.40
        184     3.40
        185     3.40
        186     3.40
        187     3.40
        188     3.40
        189     3.40
        190     3.40
        191     3.40
```

```

192    3.40
193    3.40
194    3.15
195    3.15
196    3.15
197    3.15
198    3.15
199    3.15
200    3.15
201    3.15
202    2.87
203    3.40
204    3.15
Name: stroke, Length: 205, dtype: object

```

Question #1 Answer:

Run the code below! Did you get the right code?

[Click here for the solution](#)

```

# calculate the mean vaule for "stroke" column
avg_3 = df["stroke"].astype("float").mean(axis = 0)

```

```

# replace NaN by mean value in "stroke" column
df["stroke"].replace(np.nan, avg_3, inplace = True)

```

Calculate the mean value for the 'horsepower' column:

```

In [14]: avg_4=df['horsepower'].astype('float').mean(axis=0)

```

Replace "NaN" by mean value :

```

In [15]: df['horsepower'].replace(np.nan, avg_4, inplace= True)

```

Calculate the mean value for 'peak-rpm' column:

```

In [16]: avg_5=df['peak-rpm'].astype('float').mean(axis=0)

```

Replace NaN by mean value:

```

In [17]: df['peak-rpm'].replace(np.nan, avg_5, inplace= True)

```

To see which values are present in a particular column, we can use the ".value_counts()" method:

```

In [18]: df['num-of-doors'].value_counts()

```

```

Out[18]: four      114
         two       89
         Name: num-of-doors, dtype: int64

```

We can see that four doors are the most common type. We can also use the ".idxmax()" method to automatically calculate the most common type:

```
In [19]: df['num-of-doors'].value_counts().idxmax()
```

```
Out[19]: 'four'
```

The replacement procedure is very similar to what we have seen previously:

```
In [20]: #replace the missing 'num-of-doors' values by the most frequent
df["num-of-doors"].replace(np.nan, "four", inplace = True)
```

Finally, let's drop all rows that do not have price data:

```
In [21]: # simply drop whole row with NaN in "price" column
df.dropna(subset=["price"], axis=0, inplace = True)
```

```
# reset index, because we dropped two rows
df.reset_index(drop = True, inplace = True)
```

```
In [22]: df.head()
```

```
Out[22]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	3	122	alfa-romero	gas	std	two	
1	3	122	alfa-romero	gas	std	two	
2	1	122	alfa-romero	gas	std	two	
3	2	164	audi	gas	std	four	
4	2	164	audi	gas	std	four	

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

```
[5 rows x 26 columns]
```

Good! Now, we obtain the dataset with no missing values.

Correct data format We are almost there!

The last step in data cleaning is checking and making sure that all data is in the correct format (int, float, text or other).

In Pandas, we use

.dtype() to check the data type

.astype() to change the data type

Let's list the data types for each column:

```
In [23]: df.dtypes
```

```
Out[23]: symboling          int64
normalized-losses      object
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                  float64
height                 float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
bore                   object
stroke                 object
compression-ratio      float64
horsepower             object
peak-rpm               object
city-mpg               int64
highway-mpg            int64
price                  object
dtype: object
```

As we can see above, some columns are not of the correct data type. Numerical variables should have type 'float' or 'int', and variables with strings such as categories should have type 'object'. For example, 'bore' and 'stroke' variables are numerical values that describe the engines, so we should expect them to be of the type 'float' or 'int', however, they are shown as type 'object'. We have to convert data types into a proper format for each column using the "astype()" method.

Convert data types to proper format:

```
In [24]: df[["bore", "stroke", "price"]] = df[["bore", "stroke", "price"]].astype("float")
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
#df[["price"]] = df[["price"]].astype("float")
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
print("Done")
```

Done

Let us list the columns after the conversion:

```
In [25]: df.dtypes
```

```
Out[25]: symboling          int64
normalized-losses         int64
make                      object
fuel-type                 object
aspiration                object
num-of-doors              object
body-style                object
drive-wheels              object
engine-location           object
wheel-base               float64
length                   float64
width                     float64
height                   float64
curb-weight               int64
engine-type               object
num-of-cylinders          object
engine-size               int64
fuel-system               object
bore                      float64
stroke                   float64
compression-ratio         float64
horsepower                object
peak-rpm                  float64
city-mpg                  int64
highway-mpg               int64
price                     float64
dtype: object
```

Wonderful!

Now, we finally obtain the cleaned dataset with no missing values and all data in its proper format.

Data Standardization Data is usually collected from different agencies with different formats. (Data Standardization is also a term for a particular type of data normalization, where we subtract the mean and divide by the standard deviation.

What is Standardization?

Standardization is the process of transforming data into a common format which allows the researcher to make the meaningful comparison.

Example

Transform mpg to L/100km:

In our dataset, the fuel consumption columns "city-mpg" and "highway-mpg" are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accepts the fuel consumption with L/100km standard

We will need to apply **data transformation** to transform mpg into L/100km

The formula for unit conversion is $L/100km = 235 / mpg$

We can do many mathematical operations directly in Pandas.

```
In [26]: df.head()
```

```
Out[26]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	\
0	3	122	alfa-romero	gas	std	
1	3	122	alfa-romero	gas	std	
2	1	122	alfa-romero	gas	std	
3	2	164	audi	gas	std	
4	2	164	audi	gas	std	

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111	
1	130	mpfi	3.47	2.68	9.0	111	
2	152	mpfi	2.68	3.47	9.0	154	
3	109	mpfi	3.19	3.40	10.0	102	
4	136	mpfi	3.19	3.40	8.0	115	

	peak-rpm	city-mpg	highway-mpg	price
0	5000.0	21	27	13495.0
1	5000.0	21	27	16500.0
2	5000.0	19	26	16500.0
3	5500.0	24	30	13950.0
4	5500.0	18	22	17450.0

```
[5 rows x 26 columns]
```

```
In [27]: # transform mpg to L/100km by mathematical operation (235 divided by mpg)
df['city-L/100km'] = 235/df["city-mpg"]

# check your transformed data
df.head()
```

```

Out[27]:   symboling  normalized-losses      make fuel-type aspiration \
0         3          122  alfa-romero      gas      std
1         3          122  alfa-romero      gas      std
2         1          122  alfa-romero      gas      std
3         2          164      audi      gas      std
4         2          164      audi      gas      std

      num-of-doors  body-style drive-wheels engine-location  wheel-base \
0         two  convertible      rwd      front      88.6
1         two  convertible      rwd      front      88.6
2         two   hatchback      rwd      front      94.5
3         four    sedan      fwd      front      99.8
4         four    sedan      4wd      front      99.4

      ...      fuel-system  bore  stroke  compression-ratio horsepower \
0      ...      mpfi  3.47   2.68          9.0      111
1      ...      mpfi  3.47   2.68          9.0      111
2      ...      mpfi  2.68   3.47          9.0      154
3      ...      mpfi  3.19   3.40         10.0      102
4      ...      mpfi  3.19   3.40          8.0      115

      peak-rpm  city-mpg highway-mpg      price  city-L/100km
0   5000.0      21      27  13495.0    11.190476
1   5000.0      21      27  16500.0    11.190476
2   5000.0      19      26  16500.0    12.368421
3   5500.0      24      30  13950.0     9.791667
4   5500.0      18      22  17450.0    13.055556

[5 rows x 27 columns]

```

Question #2:

According to the example above, transform mpg to L/100km in the column of "highway-mpg", and change the name of column to "highway-L/100km":

```

In [28]: df['highway-L/100km'] = 235/df["highway-mpg"]
df.head()

```

```

Out[28]:   symboling  normalized-losses      make fuel-type aspiration \
0         3          122  alfa-romero      gas      std
1         3          122  alfa-romero      gas      std
2         1          122  alfa-romero      gas      std
3         2          164      audi      gas      std
4         2          164      audi      gas      std

      num-of-doors  body-style drive-wheels engine-location  wheel-base \
0         two  convertible      rwd      front      88.6
1         two  convertible      rwd      front      88.6
2         two   hatchback      rwd      front      94.5

```

3	four	sedan	fwd	front	99.8
4	four	sedan	4wd	front	99.4

	...	bore	stroke	compression-ratio	horsepower	peak-rpm	\
0	...	3.47	2.68	9.0	111	5000.0	
1	...	3.47	2.68	9.0	111	5000.0	
2	...	2.68	3.47	9.0	154	5000.0	
3	...	3.19	3.40	10.0	102	5500.0	
4	...	3.19	3.40	8.0	115	5500.0	

	city-mpg	highway-mpg	price	city-L/100km	highway-L/100km
0	21	27	13495.0	11.190476	8.703704
1	21	27	16500.0	11.190476	8.703704
2	19	26	16500.0	12.368421	9.038462
3	24	30	13950.0	9.791667	7.833333
4	18	22	17450.0	13.055556	10.681818

[5 rows x 28 columns]

Question #2 Answer:

Run the code below! Did you get the right code?

[Click here for the solution](#)

```
# transform mpg to L/100km by mathematical operation (235 divided by mpg)
df["highway-mpg"] = 235/df["highway-mpg"]

# rename column name from "highway-mpg" to "highway-L/100km"
df.rename(columns={"highway-mpg": "highway-L/100km"}, inplace=True)

# check your transformed data
df.head()
```

Data Normalization

Why normalization?

Normalization is the process of transforming values of several variables into a similar range. Typical normalizations include scaling the variable so the variable average is 0, scaling the variable so the variable variance is 1, or scaling the variable so the variable values range from 0 to 1.

Example

To demonstrate normalization, let's say we want to scale the columns "length", "width" and "height"

Target: We would like to Normalize those variables so their value ranges from 0 to 1.

Approach: Replace original value by (original value)/(maximum value)

```
In [29]: # replace (original value) by (original value)/(maximum value)
df['length'] = df['length']/df['length'].max()
df['width'] = df['width']/df['width'].max()
```

Question #3:

According to the example above, normalize the column "height":


```
In [30]: df['height']=df['height']/df['height'].max()
         df[["length","width","height"]].head()
```

```
Out[30]:
```

	length	width	height
0	0.811148	0.890278	0.816054
1	0.811148	0.890278	0.816054
2	0.822681	0.909722	0.876254
3	0.848630	0.919444	0.908027
4	0.848630	0.922222	0.908027

Question #3 Answer:

Run the code below! Did you get the right code?

[Click here for the solution](#)

```
df['height'] = df['height']/df['height'].max()
# show the scaled columns
df[["length","width","height"]].head()
```

Here we can see we've normalized "length", "width", and "height" in the range of [0,1].

Binning **Why binning?**

Binning is a process of transforming continuous numerical variables into discrete categorical 'bins' for grouped analysis.

Example:

In our dataset, "horsepower" is a real valued variable ranging from 48 to 288, and it has 57 unique values. What if we only care about the price difference between cars with high horsepower, medium horsepower, and little horsepower (3 types)? Can we rearrange them into three 'bins' to simplify analysis?

We will use the Pandas method 'cut' to segment the 'horsepower' column into 3 bins.

1.4 Example of Binning Data In Pandas

Convert data to correct format:

```
In [31]: df["horsepower"]=df["horsepower"].astype(float, copy=True)
```

We would like four bins of equal size bandwidth. The fourth is because the function "cut" includes the rightmost value:

```
In [32]: binwidth = (max(df["horsepower"])-min(df["horsepower"]))//4
```

We build a bin array with a minimum value to a maximum value, with bandwidth calculated above. The bins will be values used to determine when one bin ends and another begins:

```
In [33]: bins = np.arange(min(df["horsepower"]), max(df["horsepower"]), binwidth)
         bins
```

```
Out[33]: array([ 48. , 101.5, 155. , 208.5])
```

We set group names:

```
In [34]: group_names = ['Low', 'Medium', 'High']
```

We apply the function "cut" to determine what each value of "df['horsepower']" belongs to.

```
In [35]: df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names,include_low=False)
df[['horsepower', 'horsepower-binned']].head(20)
```

```
Out[35]:
```

	horsepower	horsepower-binned
0	111.0	Medium
1	111.0	Medium
2	154.0	Medium
3	102.0	Medium
4	115.0	Medium
5	110.0	Medium
6	110.0	Medium
7	110.0	Medium
8	140.0	Medium
9	101.0	Low
10	101.0	Low
11	121.0	Medium
12	121.0	Medium
13	121.0	Medium
14	182.0	High
15	182.0	High
16	182.0	High
17	48.0	Low
18	70.0	Low
19	70.0	Low

Check the dataframe above carefully, you will find the last column provides the bins for "horsepower" with 3 categories ("Low", "Medium" and "High").

We successfully narrow the intervals from 57 to 3!

1.5 Bins visualization

Normally, a histogram is used to visualize the distribution of bins we created above.


```
In [36]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot

a = (0,1,2)

# draw histogram of attribute "horsepower" with bins = 3
plt.pyplot.hist(df["horsepower"], bins = 3)

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
Out[36]: Text(0.5, 1.0, 'horsepower bins')
```



The plot above shows the binning result for attribute "horsepower".

Indicator variable (or dummy variable) **What is an indicator variable?**

An indicator variable (or dummy variable) is a numerical variable used to label categories. They are called 'dummies' because the numbers themselves don't have inherent meaning.

Why do we use indicator variables?

So that we can use categorical variables for regression analysis in the later modules.

Example

We see the column "fuel-type" has two unique values, "gas" or "diesel". Regression doesn't understand words, only numbers. To use this attribute in regression analysis, we convert "fuel-type" into indicator variables.

We will use the panda's method 'get_dummies' to assign numerical values to different categories of fuel type.

```
In [37]: df.columns
```

```
Out[37]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',  
              'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',  
              'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',  
              'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',  
              'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',  
              'highway-mpg', 'price', 'city-L/100km', 'highway-L/100km',  
              'horsepower-binned'],  
              dtype='object')
```

Get indicator variables and assign it to data frame "dummy_variable_1":

```
In [38]: dummy_variable_1 = pd.get_dummies(df["fuel-type"])
        dummy_variable_1.head()
```

```
Out[38]:   diesel  gas
0         0    1
1         0    1
2         0    1
3         0    1
4         0    1
```

Change column names for clarity:

```
In [39]: dummy_variable_1.rename(columns={'fuel-type-diesel':'gas', 'fuel-type-gas':'diesel'})
        dummy_variable_1.head()
```

```
Out[39]:   diesel  gas
0         0    1
1         0    1
2         0    1
3         0    1
4         0    1
```

We now have the value 0 to represent "gas" and 1 to represent "diesel" in the column "fuel-type". We will now insert this column back into our original dataset.

```
In [40]: # merge data frame "df" and "dummy_variable_1"
        df = pd.concat([df, dummy_variable_1], axis=1)

        # drop original column "fuel-type" from "df"
        df.drop("fuel-type", axis = 1, inplace=True)
```

```
In [41]: df.head()
```

```
Out[41]:   symboling  normalized-losses   make aspiration num-of-doors \
0         3         122  alfa-romero      std         two
1         3         122  alfa-romero      std         two
2         1         122  alfa-romero      std         two
3         2         164      audi      std         four
4         2         164      audi      std         four

   body-style drive-wheels engine-location  wheel-base  length ... \
0  convertible      rwd      front      88.6  0.811148 ...
1  convertible      rwd      front      88.6  0.811148 ...
2   hatchback      rwd      front      94.5  0.822681 ...
3      sedan      fwd      front      99.8  0.848630 ...
4      sedan      4wd      front      99.4  0.848630 ...
```

	horsepower	peak-rpm	city-mpg	highway-mpg	price	city-L/100km \
0	111.0	5000.0	21	27	13495.0	11.190476
1	111.0	5000.0	21	27	16500.0	11.190476
2	154.0	5000.0	19	26	16500.0	12.368421
3	102.0	5500.0	24	30	13950.0	9.791667
4	115.0	5500.0	18	22	17450.0	13.055556

	highway-L/100km	horsepower-binned	diesel	gas
0	8.703704	Medium	0	1
1	8.703704	Medium	0	1
2	9.038462	Medium	0	1
3	7.833333	Medium	0	1
4	10.681818	Medium	0	1

[5 rows x 30 columns]

The last two columns are now the indicator variable representation of the fuel-type variable. It's all 0s and 1s now.

Question #4:

As above, create indicator variable to the column of "aspiration": "std" to 0, while "turbo" to 1:

```
In [42]: dummy_variable_2 = pd.get_dummies(df["aspiration"])
dummy_variable_2.head()
```

```
Out[42]:
```

	std	turbo
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

```
In [43]: dummy_variable_2.rename(columns={'aspiration': 'std', 'aspiration': 'turbo'}, inplace=True)
dummy_variable_2.head()
```

```
Out[43]:
```

	std	turbo
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

```
In [44]: # merge data frame "df" and "dummy_variable_2"
df = pd.concat([df, dummy_variable_2], axis=1)

# drop original column "aspiration" from "df"
df.drop("aspiration", axis = 1, inplace=True)

df.head()
```

```

Out[44]:      symboling  normalized-losses      make num-of-doors  body-style \
0           3           122  alfa-romero      two  convertible
1           3           122  alfa-romero      two  convertible
2           1           122  alfa-romero      two   hatchback
3           2           164      audi      four      sedan
4           2           164      audi      four      sedan

      drive-wheels engine-location  wheel-base  length  width  ... \
0           rwd      front      88.6  0.811148  0.890278  ...
1           rwd      front      88.6  0.811148  0.890278  ...
2           rwd      front      94.5  0.822681  0.909722  ...
3           fwd      front      99.8  0.848630  0.919444  ...
4           4wd      front      99.4  0.848630  0.922222  ...

      city-mpg  highway-mpg  price  city-L/100km  highway-L/100km  \
0           21           27  13495.0    11.190476      8.703704
1           21           27  16500.0    11.190476      8.703704
2           19           26  16500.0    12.368421      9.038462
3           24           30  13950.0     9.791667      7.833333
4           18           22  17450.0    13.055556     10.681818

      horsepower-binned  diesel  gas  std  turbo
0           Medium      0     1     1     0
1           Medium      0     1     1     0
2           Medium      0     1     1     0
3           Medium      0     1     1     0
4           Medium      0     1     1     0

[5 rows x 31 columns]

```

Question #4 Answer:

Run the code below! Did you get the right code?

[Click here for the solution](#)

```

# get indicator variables of aspiration and assign it to data frame "dummy_variable_2"
dummy_variable_2 = pd.get_dummies(df['aspiration'])

# change column names for clarity
dummy_variable_2.rename(columns={'std': 'aspiration-std', 'turbo': 'aspiration-turbo'}, inplace=True)

# show first 5 instances of data frame "dummy_variable_1"
dummy_variable_2.head()

```

Question #5:

Merge the new dataframe to the original dataframe then drop the column 'aspiration':

In []:

Question #5 Answer:

Run the code below! Did you get the right code?
[Click here for the solution](#)

```
#merge the new dataframe to the original datafram
df = pd.concat([df, dummy_variable_2], axis=1)

# drop original column "aspiration" from "df"
df.drop('aspiration', axis = 1, inplace=True)

save the new csv

In [45]: df.to_csv('clean_df.csv')
```

2 About the Authors:

This notebook written by [Mahdi Noorian PhD](#), [Joseph Santarcangelo PhD](#), Bahare Talayian, Eric Xiao, Steven Dong, Parizad, Hima Vsudevan and [Fiorella Wenver](#). Copyright 1 2017 [cognitive-class.ai](#). This notebook and its source code are released under the terms of the [MIT License](#).

[Link](#)